

The RabbitMQ Message Broker

Martin Toshev



Agenda

- Messaging Basics
- RabbitMQ Overview
- Messaging Patterns
- Administration



Agenda

- Scalability and High Availability
- Integrations
- Security



Messaging Basics



Messaging Basics

- Messaging provides a mechanism for loosely-coupled integration of systems
- The central unit of processing in a message is a message which typically contains a **body** and a **header**



Messaging Basics

- Use cases include:
 - Log aggregation between systems
 - Event propagation between systems
 - Offloading long-running tasks to worker nodes



Messaging Basics

- Messaging solutions implement different protocols for transferring of messages such as AMQP, XMPP, MQTT and many others
- The variety of protocols implies vendor lock-in when using a particular messaging solution (also called a messaging broker)



Messaging Basics

- A variety of messaging brokers can be a choice for applications ...



Messaging Basics

- Messaging solutions provide means for:
 - securing message transfer, authenticating and authorizing messaging endpoints
 - routing messages between endpoints
 - subscribing to the broker



Messaging Basics

- An **enterprise service bus** (ESB) is one layer of abstraction above a messaging solution that further provides:
 - adapters for different for messaging protocols
 - translation of messages between the different types of protocols



RabbitMQ Overview



RabbitMQ Overview

- An open source message broker written in Erlang
- Implements the AMQP Protocol (Advanced Message Queueing Protocol)
- Has a pluggable architecture and provides extension for other protocols such as HTTP, STOMP and MQTT



RabbitMQ Overview

- AMQP is a binary protocol that aims to standardize middleware communication
- The AMQP protocol derives its origins from the financial industry - processing of large volumes of financial data between different systems is a classic use case of messaging



RabbitMQ Overview

- The AMQP protocol defines:
 - **exchanges** – the message broker endpoints that receive messages
 - **queues** – the message broker endpoints that store messages from exchanges and are used by subscribers for retrieval of messages
 - **bindings** – rules that bind exchanges and queues
- The AMQP protocol is programmable – which means that the above entities can be created/modified/deleted by applications




RabbitMQ Overview

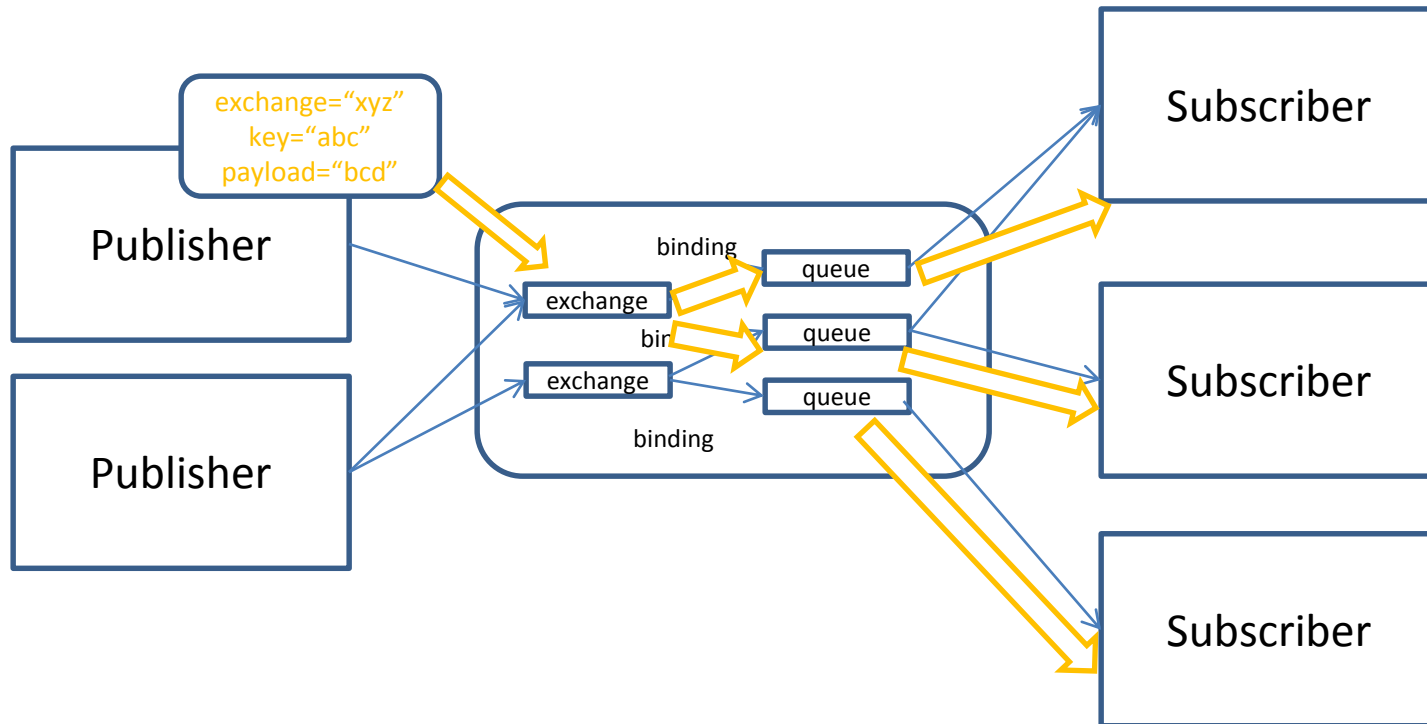
- The AMQP protocol defines multiple connection channels inside a single TCP connection in order to remove the overhead of opening a large number of TCP connections to the message broker



RabbitMQ Overview



RabbitMQ Overview



RabbitMQ Overview

- Each message can be published with a **routing key**
- Each binding between an exchange and a queue has a **binding key**
- Routing of messages is determined based on matching between the routing and binding keys



Messaging Patterns in RabbitMQ



Messaging Patterns with RabbitMQ

- Different types of messaging patterns are implemented by means of different types of exchanges
- RabbitMQ provides the following types of exchanges:
 - default
 - direct
 - fanout
 - topic
 - headers

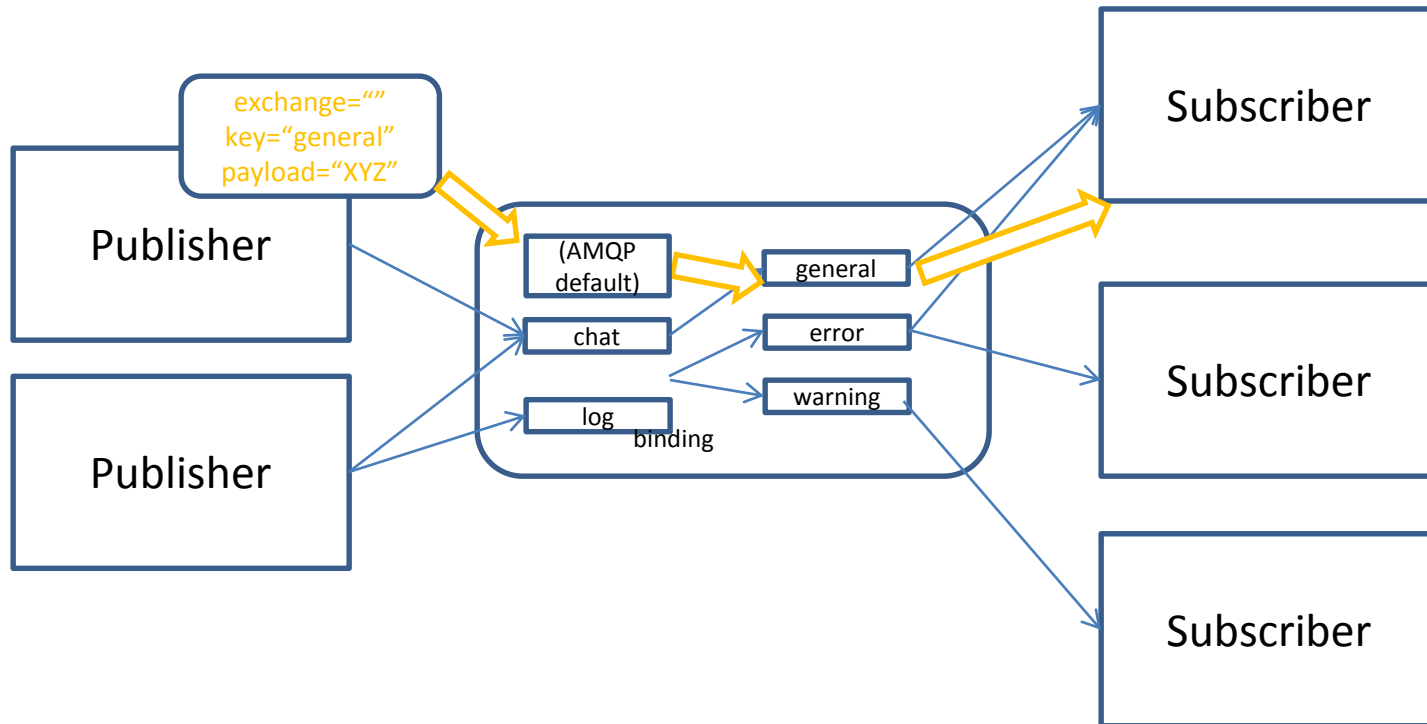


Messaging Patterns with RabbitMQ

- A default exchange has the empty string as a name and routes messages to a queue if the routing key of the message matches the queue name (no binding needs to be declared between a default exchange and a queue)
- Default exchanges are suitable for point-to-point communication between endpoints



RabbitMQ Overview



(AMQP default) is a system exchange

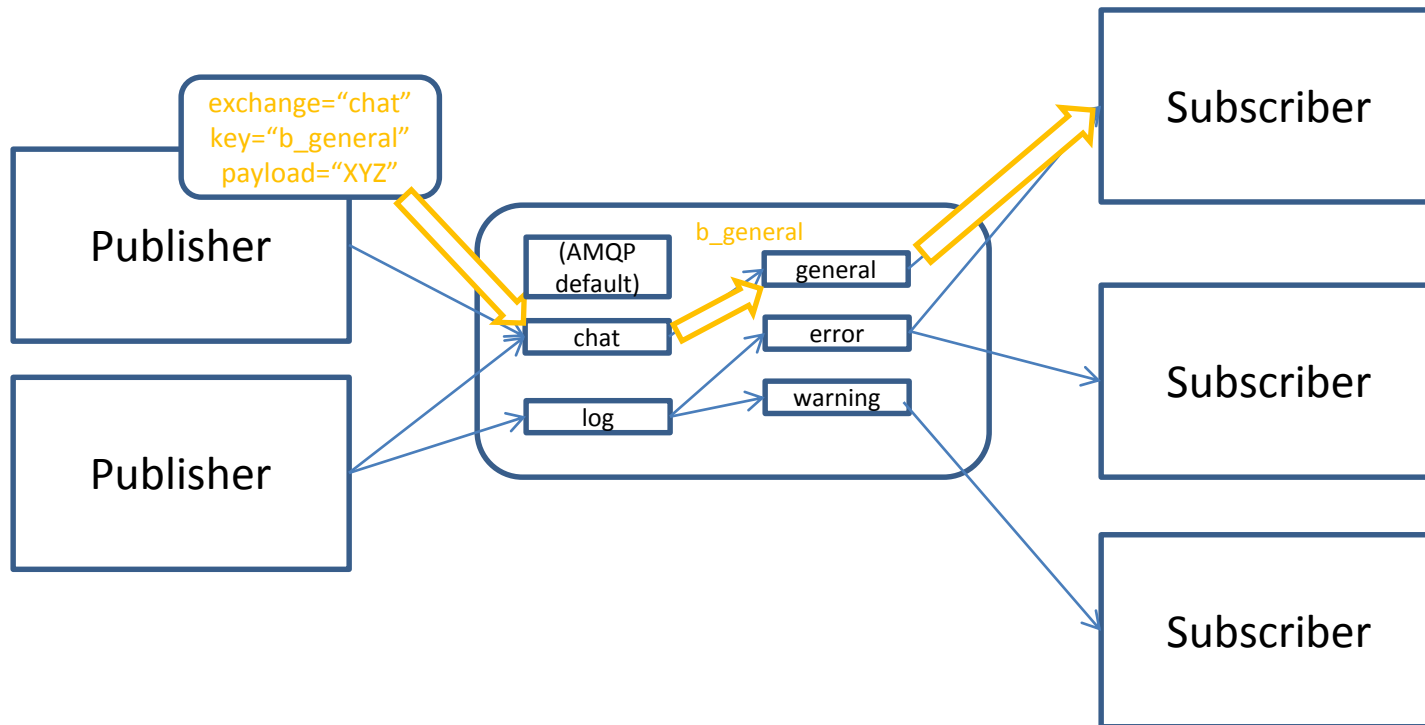


Messaging Patterns with RabbitMQ

- A direct exchange routes messages to a queue if the routing key of the message matches the binding key between the direct exchange and the queue
- Direct exchanges are suitable for point-to-point communication between endpoints



RabbitMQ Overview



***chat** is defined as a direct exchange upon creation*

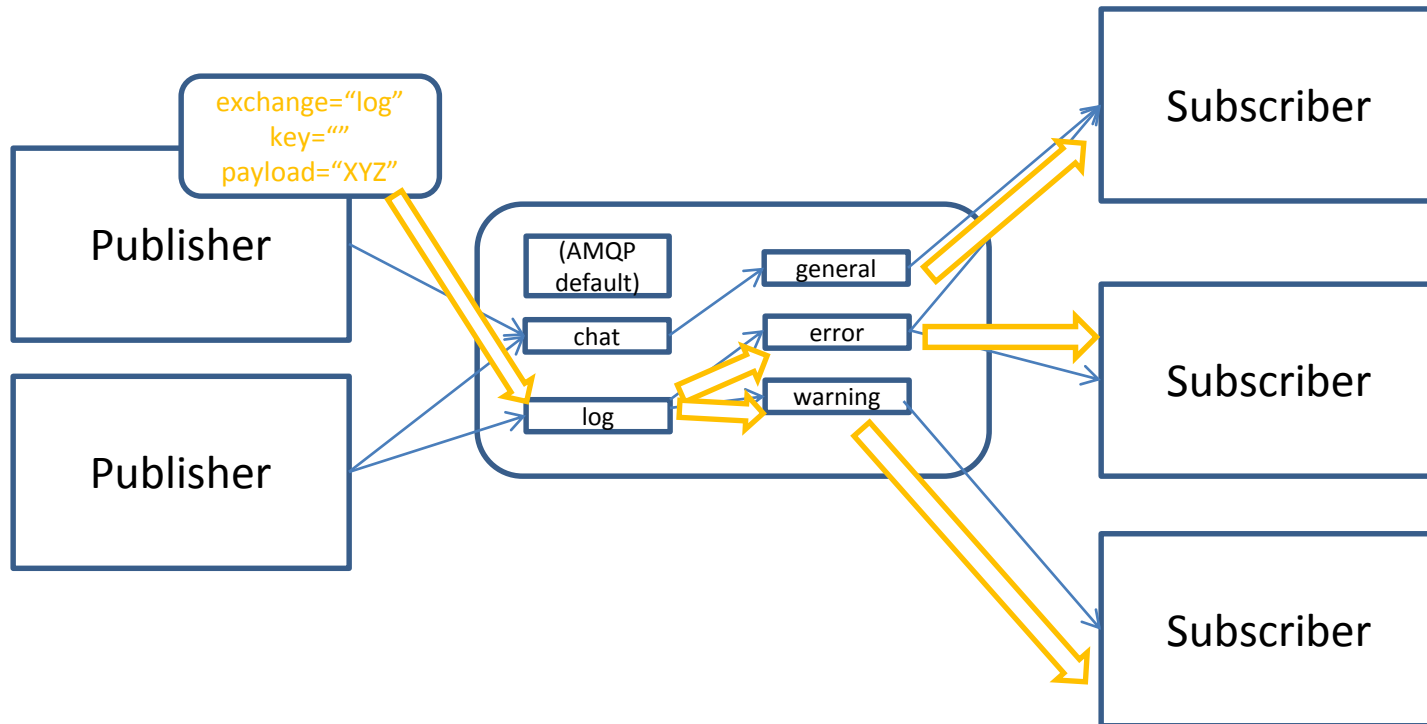


Messaging Patterns with RabbitMQ

- A fanout exchange routes (broadcasts) messages to all queues that are bound to it (the binding key is not used)
- Fanout exchanges are suitable for publish-subscribe communication between endpoints



RabbitMQ Overview



***log** is defined as a fanout exchange upon creation*

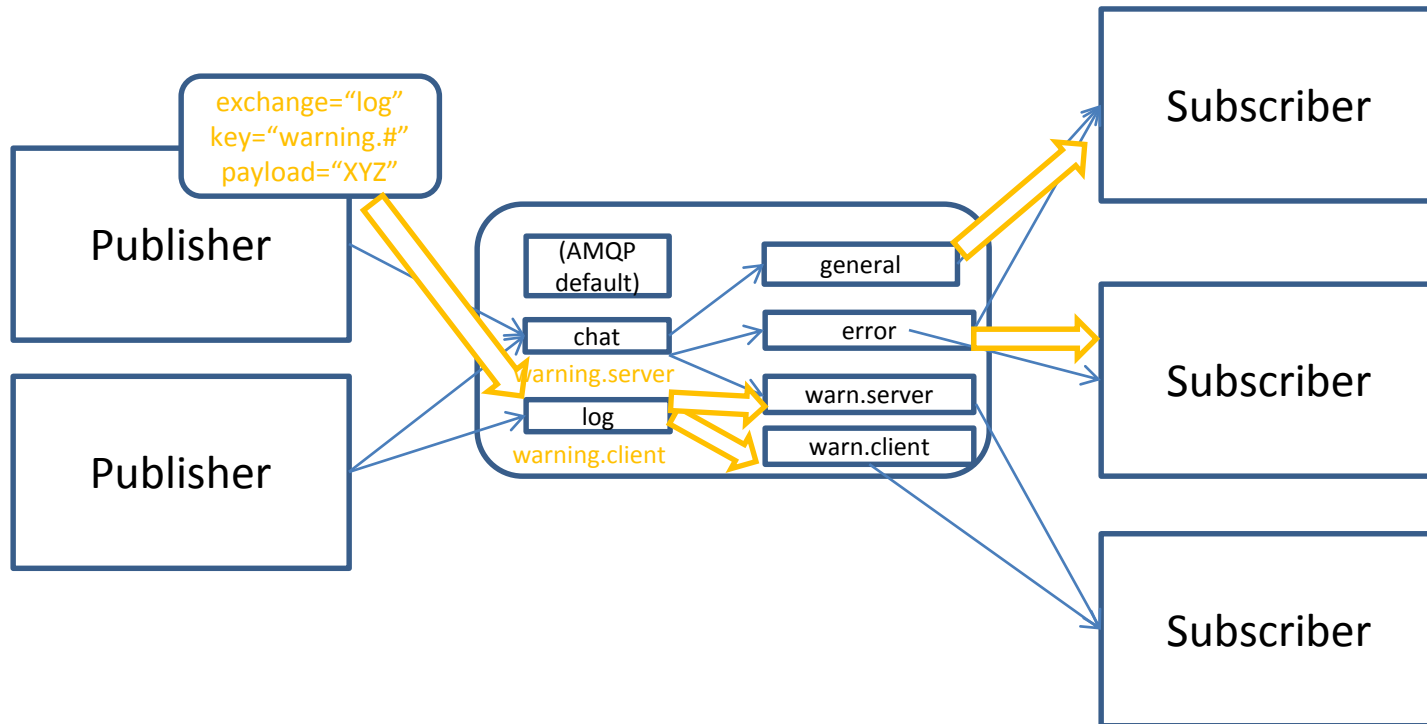


Messaging Patterns with RabbitMQ

- A topic exchange routes (multicasts) messages to all queues that have a binding key (can be a pattern) that matches the routing key of the message
- Topic exchanges are suitable for routing messages to different queues based on the type of message



RabbitMQ Overview



***log** is defined as a topic exchange upon creation*



Messaging Patterns with RabbitMQ

- A headers exchange routes messages based on a custom message header
- Header exchanges are suitable for routing messages to different queues based on more than one attribute



Messaging Patterns in RabbitMQ

(demo)



Administration



Administration

- Administration of the broker includes a number of activities such as:
 - updating the broker
 - backing up the broker database
 - Installing/uninstalling and configuring plug-ins
 - configuring the various components of the broker



Administration

- Apart from queues, exchanges and bindings we can also manage the following types of components:
 - vhosts – for logical separation of broker components
 - users
 - parameters (e.g. for defining upstream links to another brokers)
 - policies (e.g. for queue mirroring)



Administration

- Administration of single instance or an entire cluster can be performed in several ways:
 - using the management Web interface
 - using the management HTTP API
 - using the **rabbitmq-admin.py** script
 - using the **rabbitmqctl** utility



Administration

(demo)



Scalability and High Availability in RabbitMQ

- RabbitMQ provides clustering support that allows new RabbitMQ nodes to be added on the fly
- Clustering by default does not guarantee that message loss may not occur




Scalability and High Availability in RabbitMQ

- Nodes in a RabbitMQ cluster can be:
 - DISK – data is persisted in the node database
 - RAM – data is buffered only in-memory
- Nodes share only broker metadata – messages are not replicated among nodes



Scalability and High Availability in RabbitMQ

- Let's create the following three-node cluster:



Scalability and High Availability in RabbitMQ

- InstanceA node (root node):

```
set RABBITMQ_NODENAME=instanceA &  
set RABBITMQ_NODE_PORT=5770 &  
set RABBITMQ_SERVER_START_ARGS=  
    -rabbitmq_management listener [{port,33333}] &  
rabbitmq-server.bat -detached
```



Scalability and High Availability in RabbitMQ

- InstanceB node (DISK node):

```
set RABBITMQ_NODENAME=instanceB &  
set RABBITMQ_NODE_PORT=5771 &  
rabbitmq-server.bat -detached  
rabbitmqctl.bat -n instanceB stop_app  
rabbitmqctl.bat -n instanceB join_cluster instanceA@MARTIN  
rabbitmqctl.bat -n instanceB start_app
```



Scalability and High Availability in RabbitMQ

- InstanceC node (RAM node):

```
set RABBITMQ_NODENAME=instanceC &  
set RABBITMQ_NODE_PORT=5772 &  
rabbitmq-server.bat -detached  
rabbitmqctl.bat -n instanceC stop_app  
rabbitmqctl.bat -n instanceC join_cluster -ram instanceA@MARTIN  
rabbitmqctl.bat -n instanceC start_app
```



Scalability and High Availability in RabbitMQ

- However ...
- If a node that hosts a queue buffers unprocessed messages goes down – the messages are lost



Scalability and High Availability in RabbitMQ

- Default clustering mechanism provides scalability in terms of queues rather than high availability
- Mirrored queues are an extension to the default clustering mechanism that can be used to establish high availability at the broker level



Scalability and High Availability in RabbitMQ

- Mirrored queues provide queue replication over different nodes that allows a message to survive node failure
- Queue mirroring is establishing by means of a mirroring policy that specifies:
 - number of nodes to use for queue replication
 - particular nodes designated by name for queue replication
 - all nodes for queue replication



Scalability and High Availability in RabbitMQ

- The node where the queue is created is the master node – all others are slaves
- A new master node can be promoted in case the original one goes down
- A slave node is promoted to the new master in case it has fully synchronized with the old master (by default)



Scalability and High Availability in RabbitMQ

- Let's define the **test** queue in the cluster and mirror it over all other nodes:

```
rabbitmqadmin.py -N instanceA declare queue name=test  
durable=false  
rabbitmqctl -n instanceA set_policy ha-all "test" '{"ha-  
mode":"" : "all"}'
```



Scalability and High Availability in RabbitMQ

- However ...
- The RabbitMQ clustering mechanism uses Erlang message passing along with a message cookie in order to establish communication between the nodes
- ... which is not reliable over the WAN ...




Scalability and High Availability in RabbitMQ

- In order to establish high availability among nodes in different geographic locations you can use the federation and shovel plug-ins
- The shovel plug-ins works at a lower level than the federation plug-in



Scalability and High Availability in RabbitMQ

- Assuming we have a remote node **instanceD** we can make the **test** queue federated on that node:



Scalability and High Availability in RabbitMQ

- Declare the remote **instanceD** instance and enable the federation plug-in for it:

```
set RABBITMQ_NODENAME=instanceD &
set RABBITMQ_NODE_PORT=6001 &
set RABBITMQ_SERVER_START_ARGS=
    -rabbitmq_management listener [{port,44444}] &
rabbitmq-server.bat -detached

rabbitmq-plugins -n instanceD enable rabbitmq_federation
rabbitmq-plugins -n instanceD enable
rabbitmq_federation_management
```



Scalability and High Availability in RabbitMQ

- Declare the **federated_test** queue:

```
rabbitmqadmin.py -N instanceD -P 44444 declare queue  
name=federated_test durable=false
```



Scalability and High Availability in RabbitMQ

- Declare the upstream to the initial cluster and set a federation link to the **test** queue:

```
rabbitmqctl -n instanceD set_parameter federation-upstream  
upstream  
"{\"uri\":\"amqp://localhost:5770\",\"expires\":3600000,  
\"queue\":\"test\"}"  
  
rabbitmqctl -n instanceD set_policy federate-queue  
--apply-to queues "federated_test"  
"{\"federation-upstream\":\"upstream\"}"
```



Scalability and High Availability in RabbitMQ

- The shovel plug-in provides two variants:
 - **static** (all links between the source/destination nodes/clusters) are defined statically in the RabbitMQ configuration file
 - **dynamic** (all links between the source/destination nodes/clusters) are defined dynamically via RabbitMQ parameters



Scalability and High Availability in RabbitMQ

- The shovel plug-in provides two variants:
 - **static** (all links between the source/destination nodes/clusters) are defined statically in the RabbitMQ configuration file
 - **dynamic** (all links between the source/destination nodes/clusters) are defined dynamically via RabbitMQ parameters



Scalability and High Availability in RabbitMQ

destination source	exchange	queue
exchange	federation dynamic shovel	dynamic shovel
queue	static shovel dynamic shovel	federation dynamic shovel



Scalability and High Availability in RabbitMQ (demo)



Integrations



Integrations

- RabbitMQ provides integrations with other protocols such as STOMP, MQTT and LDAP by means of RabbitMQ plug-ins
- The Spring Framework provides integration with AMQP protocol and RabbitMQ in particular



Integrations

- The Spring AMQP framework provides:
 - **RabbitAdmin** class for automatically declaring queues, exchanges and bindings
 - Listener container for asynchronous processing of inbound messages
 - **RabbitTemplate** class for sending and receiving messages



Integrations

- Utilities of the Spring AMQP framework can be used either directly in Java or preconfigured in the Spring configuration
- The Spring Integration framework provides adapters for the AMQP protocol



Integrations

- The following Maven dependency can be used to provide the Spring AMQP framework to the application:

```
<dependencies>
  <dependency>
    <groupId>org.springframework.amqp</groupId>
    <artifactId>spring-rabbit</artifactId>
    <version>1.4.5.RELEASE</version>
  </dependency>
</dependencies>
```



Integrations

- **RabbitAdmin** (plain Java):

```
CachingConnectionFactory factory = new
    CachingConnectionFactory("localhost");
RabbitAdmin admin = new RabbitAdmin(factory);
Queue queue = new Queue("sample-queue");
admin.declareQueue(queue);
TopicExchange exchange = new TopicExchange("sample-topic-
    exchange");
admin.declareExchange(exchange);
admin.declareBinding(BindingBuilder.bind(queue).to(exchange)
    .with("sample-key"));
factory.destroy();
```



Integrations

- Container listener (plain Java):

```
CachingConnectionFactory factory =  
    new CachingConnectionFactory(  
        "localhost");  
SimpleMessageListenerContainer container = new  
SimpleMessageListenerContainer(  
    factory);  
Object listener = new Object() {  
    public void handleMessage(String message) {  
        System.out.println("Message received: " +  
            message);  
    }  
};  
MessageListenerAdapter adapter = new  
    MessageListenerAdapter(listener);  
container.setMessageListener(adapter);  
container.setQueueNames("sample-queue");  
container.start();
```



Integrations

- **RabbitTemplate** (plain Java):

```
CachingConnectionFactory factory =  
    new CachingConnectionFactory("localhost");  
RabbitTemplate template = new  
RabbitTemplate(factory);  
template.convertAndSend("", "sample-queue",  
    "sample-queue test message!");
```



Integrations

- All of the above examples can be configured using the Spring configuration
- Must cleaner and decouples RabbitMQ configuration for the business logic



Integrations

(demo)



Security



Security

- RabbitMQ uses SASL for authentication (SASL PLAIN used by default)
- RabbitMQ uses access control lists (permissions) for authorization



Security

- SSL/TLS support can be enabled for the AMQP communication channels
- SSL/TLS support can be enabled for node communication between nodes in a cluster
- SSL/TLS support can be enabled for the federation and shovel plug-ins




Summary

- RabbitMQ is a robust messaging solution that can be used in a variety of scenarios based on your application needs
- RabbitMQ may not be the best possible solution compared to other messaging brokers – always consider benchmarks based on size and number of messages



Readings



Thanks

?

